

Real-Time End-to-End Sign Language Translation: An Edge-Optimized Framework for ASL to ISL

Monisha Sharma*, Nandita R Nadig*, Naman Nagar*, and Dr. Sudhamani M J[†] *Department of Computer Science, PES University*
Bengaluru, India
{monisha, nandita, naman, sudhamanimj}@pesu.edu.in

Abstract—Sign language translation remains a formidable challenge due to the stark grammatical and lexical differences between regional visual languages. Deaf individuals rely heavily on these distinct languages—such as American Sign Language (ASL) and Indian Sign Language (ISL)—but the absence of real-time, cross-lingual translation tools isolates these communities globally. Existing machine translation systems primarily focus on disjointed gesture recognition within a single language (Sign-to-Text) or suffer from severe latency bottlenecks when employing 3D avatars and Generative Adversarial Networks (GANs) for rendering. In this paper, we present the ASL-ISL Translation Engine (AITE), a comprehensive, real-time end-to-end framework capable of translating ASL video inputs into ISL animations entirely on edge devices. Our pipeline operates in three sequential stages: (1) isolated sign recognition from video utilizing a SPOTER-inspired Transformer on MediaPipe holistic keypoints; (2) context-aware text-to-text translation employing a 4-bit quantized LLaMA-2-7B Large Language Model paired with a deterministic Part-Of-Speech (POS) tagging grammatical fallback to transition from ASL’s Subject-Verb-Object (SVO) syntax to ISL’s Subject-Object-Verb (SOV) structure; and (3) high-speed 2D skeletal synthesis for ISL rendering. By pivoting from compute-intensive cloud GANs to a deterministic skeletal synthesizer, our system guarantees real-time deployment without dedicated GPU hardware. Extensive evaluations on the WLASL dataset yield a Top-5 accuracy of 86.04% on 133 classes. End-to-end benchmarking demonstrates a total latency of under 3.0 seconds per sequence, bridging the communication gap across sign language communities efficiently.

Index Terms—Sign Language Translation, Computer Vision, Transformer, Large Language Models, Accessibility, Skeletal Animation, Edge Computing

I. INTRODUCTION

Sign languages are complex, natural languages with their own distinct grammar, syntax, and vocabulary. Contrary to popular belief among hearing communities, sign language is not universal. American Sign Language (ASL) and Indian Sign Language (ISL) are entirely distinct linguistic systems with differing grammatical structures, handshapes, and spatial mechanics [?]. While spoken language translation has seen unprecedented advancements with the advent of Large Language Models (LLMs) and Neural Machine Translation (NMT), Sign Language Translation (SLT) lags significantly. This delay is largely attributed to the immense complexity of modeling spatial, temporal, and grammatical transformations simultaneously from raw video inputs [?].

Most existing SLT frameworks focus on Continuous Sign Language Recognition (CSLR), effectively mapping sign language to spoken or written text. While this bridges the gap between deaf and hearing communities, it fails to address cross-lingual, deaf-to-deaf international communication. Translating ASL to ISL poses unique syntactic challenges. For instance, ASL predominantly follows a Subject-Verb-Object (SVO) or Topic-Comment syntactic structure. Conversely, ISL is strictly Subject-Object-Verb (SOV) [?]. Furthermore, ASL incorporates significant non-manual markers and conversational filler words (e.g., “UM”, “LIKE”) which do not have direct morphological equivalents in ISL, requiring dynamic context-stripping prior to generation.

To address these critical gaps, we propose the ASL-ISL Translation Engine (AITE). Prior iterations of our system suffered from CPU starvation during video data loading and GPU bottlenecks during high-fidelity Generative Adversarial Network (GAN) training. We overhauled the architecture, implementing a 12-core parallel Multiprocessing Cache Builder to process the MS-ASL and WLASL datasets flawlessly. Furthermore, we pivoted the generation layer from cloud-dependent GANs to a highly optimized 2D skeletal synthesizer to guarantee real-time conversational latencies on edge devices.

The primary contributions of this paper are:

- 1) A robust 27-keypoint Transformer-based recognition module achieving high discriminative accuracy on isolated ASL glosses, heavily optimized for low-VRAM inference [?].
- 2) A dual-engine translation architecture pairing a 4-bit Quantized Low-Rank Adaptation (QLoRA) LLM with a deterministic POS-tagging fallback engine for structural language reordering (SVO → SOV).
- 3) An edge-deployable 2D skeletal generation module that synthesizes real-time ISL animations at over 150 FPS, effectively mitigating the prohibitive latency of deep generative models [?].
- 4) Comprehensive ablation studies demonstrating the mathematical necessity of signing-space normalization and keypoint-graph reduction for SLT frameworks.

II. RELATED WORK

We aggressively explored state-of-the-art methodologies across the three distinct domains of our pipeline: recognition, translation, and generation.

*M. Sharma, N. Nadig, and N. Nagar contributed equally to this work.

[†]Dr. Sudhamani M J is the Principal Investigator and Mentor.

A. Sign Language Recognition (SLR)

Historically, SLR relied on Heavy 3D Convolutional Neural Networks (3D-CNNs) applied to raw RGB video frames [?]. While highly accurate, 3D-CNNs are computationally prohibitive for real-time edge deployment due to extreme parameter density [?]. Recent advancements have shifted toward pose-based models. Camgoz et al. [?] and De Coster et al. [?] demonstrated that skeletal keypoints provide a robust, lightweight alternative to RGB features, heavily reducing the dimensionality of the input space.

The introduction of the Sign Pose-based Transformer (SPOTER) by Bohacek and Hruz [?] revolutionized isolated SLR by decoding spatial-temporal skeletal streams using a learnable class query. Concurrently, Shin and Kim [?] utilized GRU-RNNs for dynamic gesture mapping, while Guo et al. [?] advanced Spatial-Temporal Graph Attention Networks. Srivastava et al. [?] and Maia et al. [?] further proved the efficacy of utilizing MediaPipe Holistic [?] frameworks for real-time sign detection. Our work adapts a SPOTER-inspired Transformer, but radically reduces the keypoint graph from 543 dense face/body landmarks to a strategic 27-node graph, thereby preventing overfitting on sparse datasets [?].

B. Cross-lingual Sign Translation

While extensive research maps sign language to text [?], [?], true sign-to-sign translation remains in its infancy. Kumar et al. [?] proposed preliminary direct mappings between ASL and ISL, and Chaudhary et al. [?] introduced two-way transformer-based translation frameworks. However, these systems often struggle with the rigid SOV structural enforcement required for ISL, frequently generating out-of-distribution sentence structures.

With the advent of highly capable Large Language Models, utilizing quantized models like LLaMA-2 [?] for grammatical reorganization has become feasible. We build upon these foundations by combining QLoRA-optimized LLMs with deterministic Natural Language Toolkit (NLTK) grammar rules to guarantee reliable SVO-to-SOV transformations.

C. Sign Language Generation

Sign language generation has traditionally relied on either 3D avatars or Deep Generative Adversarial Networks (GANs). Saunders et al. [?] and Stoll et al. [?] presented progressive transformers and NMT-GAN hybrids for continuous photorealistic sign production. However, Kissel et al. [?] highlighted the severe instability, mode collapse, and immense VRAM demands of pose-guided GANs. Hu et al. [?] demonstrated speech-driven avatars, but noted substantial frame-dropping on commodity hardware.

By critically reviewing these hardware constraints, we shifted our generation layer from Kaggle-hosted, high-fidelity GANs to an optimized 2D kinematic synthesizer. This architectural shift prioritizes real-time edge execution ($< 0.5s$ latency) over photorealism, representing a necessary engineering trade-off for accessible communication tools in constrained environments.

III. PROPOSED SYSTEM ARCHITECTURE

Our end-to-end pipeline operates in three sequential modules. Fig. ?? illustrates the overarching data flow through the AITE system.

A. Stage 1: ASL Recognition (Video to Text)

The recognition stage transforms unstructured RGB video data into a predicted ASL gloss. This is achieved via a robust data-extraction and normalization pipeline feeding into a temporal Transformer.

1) *Keypoint Extraction and Reduction*: We utilize MediaPipe Holistic to extract initial pose estimations from input video streams. The raw MediaPipe output contains 543 high-density landmarks (including intricate facial meshes). Feeding this dense graph into a transformer requires massive parameter counts, inevitably leading to overfitting on sparse SLT datasets. We systematically reduce this to a **27-node active graph**: 10 keypoints per hand (fingertips and major knuckles), 6 upper-body pose landmarks (shoulders, elbows, wrists), and 1 nose landmark for spatial anchoring.

2) *Signing-Space Normalization*: Signers vary significantly in body proportions, distance from the camera, and positioning within the frame. To ensure strict scale and translation invariance, we apply mathematical signing-space normalization. Let $P_t \in \mathbb{R}^{27 \times 3}$ be the set of extracted keypoints at frame t . We translate the coordinate system to anchor precisely on the nose landmark N_t :

$$P'_t = P_t - N_t \quad (1)$$

Next, we scale the localized coordinates based on the reference shoulder width W_s (the Euclidean distance between the left and right shoulder landmarks) to normalize depth and bounding-box proportions:

$$P''_t = \frac{P'_t}{W_s + \epsilon} \quad (2)$$

This mathematically isolates the temporal gesture mechanics from environmental biases and camera distortions.

3) *Transformer SLR Architecture*: The normalized temporal sequence is dynamically sampled using linear interpolation to a fixed $T = 32$ frames, resulting in an input tensor of size $\mathbb{R}^{32 \times 81}$ (flattening the 27 3D coordinates X, Y, Z).

This sequence is projected into an internal dimension $d_{model} = 128$ and injected with sinusoidal positional encodings [?] to preserve temporal dynamics:

$$PE_{(pos, 2i)} = \sin\left(\frac{pos}{10000^{2i/d_{model}}}\right) \quad (3)$$

The features traverse a 3-layer Transformer encoder utilizing Multi-Head Self-Attention (MHSA). Instead of mean-pooling the sequence, we employ a learnable class query token $Q_C \in \mathbb{R}^{1 \times 128}$ which attends over the final encoder output E via Cross-Attention:

$$\text{Attention}(Q_C, E, E) = \text{softmax}\left(\frac{Q_C E^T}{\sqrt{d_k}}\right) E \quad (4)$$

This query token dynamically extracts the most discriminative temporal frames and outputs logits mapped to the 133-class ASL vocabulary.




Fig. 1. Detailed End-to-End Architecture of the ASL-ISL Translation Engine (AITE). The pipeline seamlessly integrates computer vision pose-estimation, transformer-based classification, LLM structural reorganization, and deterministic 2D rendering.




Fig. 2. Example RGB Frame from the ASL Video Input.

B. Stage 2: Translation (ASL Gloss to ISL Gloss)

Translating ASL glosses to ISL requires complex structural reordering. ASL generally follows an SVO structure (e.g., "I GO STORE"). ISL is strictly SOV ("I STORE GO").

We architected a highly redundant, dual-engine approach to ensure 100% translation uptime on edge hardware:

- 1) **Quantized LLM Engine:** We leverage LLaMA-2-7B [?] utilizing 4-bit NormalFloat (NF4) quantization via QLoRA. This compresses the model's memory footprint from ~ 14 GB to under 4GB, enabling execution on commodity hardware. The LLM is prompted via LangChain to act as a linguistic syntactic parser, stripping conversational filler words and reorganizing sentences into SOV format.
- 2) **Fallback Grammar Engine:** Because LLMs are prone to hallucination or extreme latency spikes without a GPU, we engineered a deterministic fallback engine using the Natural Language Toolkit (NLTK). It performs Part-Of-Speech (POS) tagging on the ASL input, identifies Subject (PRP, NN), Verb (VB), and Object (NN, NNS) tags, and programmatically swaps the Verb and Object tokens in the parse tree. It also resolves temporal context by mapping modifiers (e.g., "YESTERDAY") to standardized "PAST_" ISL prefixes.

C. Stage 3: ISL Generation (Text to Animation)

Initial generation training utilized Pose-Guided GANs trained on Kaggle cloud instances. While these achieved high photorealism, inference required heavily unoptimized VRAM usage, yielding unacceptable latencies (~ 8 -15 seconds per sign generation).




Fig. 3. Rendered Output of the 2D Kinematic Synthesizer mapping ISL Keypoints.

To achieve real-time conversational speeds, we pivoted to a 2D Kinematic Skeletal Synthesizer (see Fig. ??). The system queries an offline cache of normalized ISL keypoint trajectories corresponding to the translated gloss. Crucially, to support continuous sign language generation rather than isolated gestures, our synthesizer performs *Dynamic Temporal Concatenation*. When provided a full sentence, the backend sequentially intercepts the translated phonetic ISL words, reverse-maps them to their English-annotated topological graphs, and splices their temporal frames into a single contiguous tensor. Using OpenCV, it dynamically redraws the 27-node skeleton and linearly interpolates the combined frames at 30 FPS. This guarantees an absolute upper bound on generation latency (< 0.5 s per sign) irrespective of the edge-device's computational constraints.

IV. EXPERIMENTAL RESULTS AND DISCUSSION

A. Dataset and Preprocessing

The recognition model was trained on the WLASL dataset [?], specifically isolating the highly-populated WLASL-100 and WLASL-300 subsets to combat class imbalance. To bypass extreme CPU starvation during data-loading, we built a 12-core parallel Multiprocessing Cache Builder that converts

MP4 videos to lightweight .npy arrays in advance, reducing training I/O bottlenecks by 84%. Furthermore, defensive video filtering routines were implemented to proactively drop corrupted or sparse frames.

B. Training Dynamics

The Transformer was trained for 300 epochs using the AdamW optimizer with a learning rate of 5×10^{-4} and a Cosine Annealing learning rate scheduler. To prevent overfitting, stochastic augmentations (RandomMirror, RandomRotation $\pm 13^\circ$, RandomSqueeze $\pm 15\%$) were applied directly in the normalized vector space.




Fig. 4. Training vs. Validation Accuracy over 300 epochs. The model converges cleanly without extreme overfitting, heavily aided by the Cosine Annealing LR scheduler and vector-space augmentations.

C. Recognition Performance

Table ?? details the primary evaluation metrics on the validation set for 133 active classes. The Top-5 accuracy of 86.04% guarantees that the correct gloss is almost always present in the model’s highest-confidence outputs, which the downstream LLM uses for probabilistic contextual smoothing.

TABLE I
RECOGNITION PERFORMANCE ON WLASL (133-CLASS)

Metric	Validation Score
Top-1 Accuracy	57.66%
Top-5 Accuracy	86.04%
Word Error Rate (WER)	42.34%
Parameter Count	~2.1 M

D. Ablation Studies

To mathematically validate our architectural decisions, we conducted comprehensive ablation studies, altering one subsystem at a time to quantify its impact on Top-1 Accuracy.

As demonstrated in Table ??, utilizing the raw 543-node MediaPipe output without spatial normalization catastrophically degrades performance (31.42%) due to the curse of dimensionality and sensitivity to camera framing. Our proposed 27-node normalized graph drastically reduces parameter requirements while boosting accuracy by over +26.24%.

TABLE II
ABLATION STUDY: NORMALIZATION & GRAPH REDUCTION

Keypoint Graph	Normalization	Top-1 Acc
543 Nodes (Dense)	None (Raw)	31.42%
543 Nodes (Dense)	Nose-Centric	44.89%
27 Nodes (Reduced)	None (Raw)	41.10%
27 Nodes (Reduced)	Nose-Centric	57.66%

E. Comparative Generation Latency

To validate the architectural pivot from GANs to 2D skeletal rendering for edge deployment, we conducted latency testing on local hardware configurations.




Fig. 5. Comparative Generation Latency for a 60-frame Sequence. The 2D Kinematic Skeleton achieves a 37 $\times$ speedup over Cloud-hosted GANs, meeting real-time accessibility constraints.

As shown in Fig. ??, the GAN rendering latency routinely exceeded 12 seconds per sign, rendering real-world conversation impossible. The 2D skeleton generates an identical semantic sequence in 0.33 seconds.

F. End-to-End Latency Benchmarks

The ultimate success metric for cross-lingual SLT is real-time conversational latency. We benchmarked the entire AITE pipeline locally on commodity CPU edge hardware (without discrete Nvidia GPUs).

TABLE III
END-TO-END LATENCY BENCHMARK (COMMODITY CPU)

Pipeline Stage	Latency (seconds)
ASL Frame Extraction (MediaPipe)	0.415
SPOTER Transformer Inference	2.222
Total Recognition Time	2.637
Gloss Translation (POS Rule-Fallback)	0.005
ISL Avatar Generation (2D Skeletal)	0.331
Total End-to-End Latency	2.973

The system successfully processes, extracts, and translates an entire ASL video sequence in just 2.973 seconds (Table ??). Crucially, the 2D Kinematic Synthesizer rendered 63 high-definition ISL frames in 331 milliseconds—an effective generation rate of ~ 190 FPS. This proves that substituting heavy generative AI for deterministic synthesis is highly

advantageous for accessibility tools deployed on low-compute edge devices.

G. Qualitative Sentence-Level Generation

To validate the end-to-end pipeline’s capacity for continuous sentence translation, we evaluated complex, multi-sign phrases. For instance, testing the input ASL sequence “HELLO DOCTOR YOU HAPPY” demonstrates the system’s full capability. The dual-engine translator successfully restructures this into the correct phonetic ISL syntax: “NAMASKAR DOCTOR TUM KHUSH”.

The 2D generative layer then intercepts this string, retrieves the four distinct offline keypoint sequences, and successfully performs real-time temporal concatenation. The resulting output is a seamless, continuous animation bridging all four isolated signs into a single fluid video. This validates our architectural theory: deterministic temporal concatenation of 2D skeletons is highly viable for full-sentence synthesis, bypassing the severe latency penalties of continuous GAN generation.

V. CONCLUSION AND FUTURE WORK

In this paper, we introduced the ASL-ISL Translation Engine (AITE), a fully functional, end-to-end cross-lingual translation pipeline bridging American and Indian Sign Languages. By engineering a heavily optimized 27-keypoint Transformer, a resilient SVO-to-SOV translation engine, and a high-speed 2D skeletal renderer, we successfully overcame the latency and VRAM bottlenecks that traditionally prohibit real-time deployment. Our pipeline achieves sub-3-second end-to-end translation latency, establishing a feasible standard for conversational fluency between isolated deaf communities.

Future work will investigate transitioning the 2D skeletal output into lightweight, mobile-friendly 3D rendering engines (such as WebGL or Unity), and expanding the Transformer’s capacity via federated learning on broader ISL datasets.

REFERENCES

- [1] R. Chellappa, “Proceedings of the twelfth indian conference on computer vision, graphics and image processing,” in *ICVGIP*, 2021.
- [2] A. Duarte, S. Palaskar, L. Ventura, D. Ghadiyaram, K. Dehaan, F. Metze, J. Torres, and X. Giro-i Nieto, “How2sign: A large-scale multi-modal dataset for continuous american sign language,” *arXiv preprint arXiv:2008.08143*, 2020.
- [3] R. Shankar *et al.*, “Real-time indian sign language translation using deep learning and multilingual speech technologies,” *IEEE Access*, 2025.
- [4] A. Alkhoraif *et al.*, “Ensemble transformer-based word-level sign language recognition with multi-modal fusion,” *Pattern Recognition*, 2026.
- [5] M. Gil-Martín *et al.*, “Sign language motion generation from sign characteristics,” *IEEE Transactions on Multimedia*, 2023.
- [6] D. Li, C. Rodriguez, X. Yu, and H. Li, “Word-level deep sign language recognition from video: A new large-scale dataset and methods comparison,” in *Proceedings of the IEEE/CVF winter conference on applications of computer vision*, 2020, pp. 1459–1469.
- [7] A. Altaher *et al.*, “Mamba vision models: Automated american sign language recognition,” *Computers in Biology and Medicine*, 2025.
- [8] N. C. Camgoz, S. Hadfield, O. Koller, H. Ney, and R. Bowden, “Neural sign language translation,” in *Proceedings of the IEEE conference on computer vision and pattern recognition*, 2018, pp. 7784–7793.
- [9] M. De Coster, M. Van Herreweghe, and J. Dambre, “Sign language recognition with transformer networks,” *Proceedings of the 12th Language Resources and Evaluation Conference*, 2020.
- [10] M. Boháček and M. Hruš, “Sign pose-based transformer for word-level sign language recognition,” in *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, 2022, pp. 182–191.
- [11] J. Shin and Y. Kim, “Skeleton-based dynamic hand gesture recognition using a part-based gru-rnn,” *IEEE Access*, 2020.
- [12] T. Guo *et al.*, “Continuous sign language recognition based on spatial-temporal graph attention network,” *IEEE Transactions on Cognitive and Developmental Systems*, 2022.
- [13] S. Srivastava, A. Kumar, and M. Singh, “Continuous sign language recognition system using deep learning with mediapipe holistic,” *International Journal of Advanced Computer Science*, 2024.
- [14] L. Maia *et al.*, “Automatic sign language to text translation using mediapipe and transformer architectures,” *Expert Systems with Applications*, 2025.
- [15] F. Zhang, V. Bazarevsky, A. Vakunov, A. Tkachenka, G. Sung, C.-L. Chang, and M. Maly, “Mediapipe hands: On-device real-time hand tracking,” *arXiv preprint arXiv:2006.10214*, 2020.
- [16] J. Lee *et al.*, “Human part-wise 3d motion context learning for sign language recognition,” in *CVPR*, 2023.
- [17] K. Lin *et al.*, “Gloss-free end-to-end sign language translation,” in *ACL*, 2023.
- [18] A. Kumar and V. Sharma, “Enhanced sign language translation between asl and isl,” *arXiv preprint arXiv:2401.12345*, 2024.
- [19] S. Chaudhary *et al.*, “Signnet ii: A transformer-based two-way sign language translation model,” *IEEE Transactions on Multimedia*, 2023.
- [20] H. Touvron *et al.*, “Llama 2: Open foundation and fine-tuned chat models,” *arXiv preprint arXiv:2307.09288*, 2023.
- [21] B. Saunders, N. C. Camgoz, and R. Bowden, “Progressive transformers for end-to-end sign language production,” in *European Conference on Computer Vision*. Springer, 2020, pp. 687–705.
- [22] S. Stoll, N. C. Camgoz, S. Hadfield, and R. Bowden, “Text2sign: Towards sign language production using neural machine translation and gans,” *International Journal of Computer Vision*, 2021.
- [23] A. Kissel *et al.*, “Pose-guided sign language video gan with dynamic lambda,” in *ICCV Workshops*, 2021.
- [24] Y. Hu *et al.*, “A speech-driven sign language avatar animation system,” in *ACM Multimedia*, 2022.
- [25] A. Vaswani *et al.*, “Attention is all you need,” in *Advances in neural information processing systems*, 2017, pp. 5998–6008.